



k8s

Kubernetes





What is Kubernetes ?

- Portable, extensible, open-source platform for managing containerized workloads and services
- Facilitates both declarative configuration and automation
- It has a large, rapidly growing ecosystem
- Kubernetes services, support, and tools are widely available
- The name Kubernetes originates from Greek, meaning helmsman or pilot
- Google open-sourced the Kubernetes project in 2014

① declarative config → YAML
↳ manifest file

② automation → CLI

K8S installation

→ on-prem → self managed → installation / updating / security by org

Cloud

→ self-managed → - - -
→ managed service → installation / update ... → responsibility of cloud provider

→ AWS → EKS → Elastic K8S service

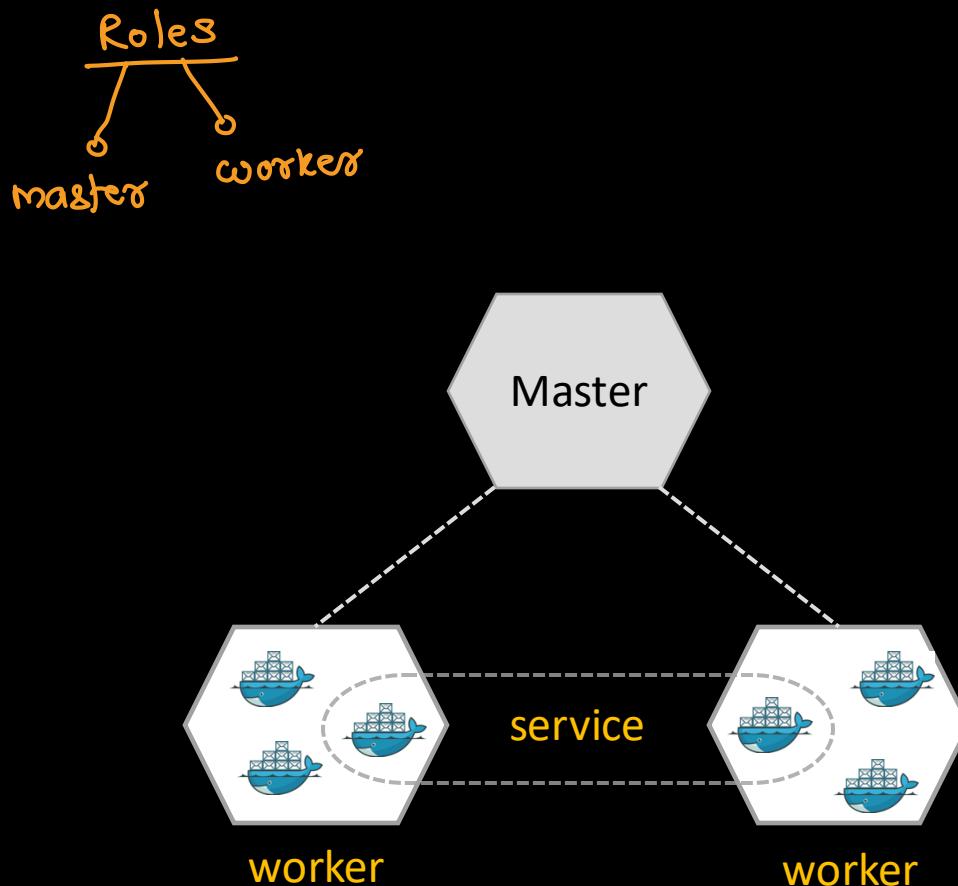
→ Azure → AKS → Azure K8S service

→ GCP → GKE → GCP K8S Engine



Kubernetes Cluster

- When you deploy Kubernetes, you get a cluster.
- A cluster is a set of machines (nodes), that run containerized applications managed by Kubernetes
- A cluster has at least one worker node and at least one master node
- The worker node(s) host the pods that are the components of the application
- The master node(s) manages the worker nodes and the pods in the cluster
- Multiple master nodes are used to provide a cluster with failover and high availability → production env



K8S clusters

→ Single Node cluster

- simulated / virtual cluster
- configured by tool → minikube
- used only for learning

→ multi node cluster

→ single master cluster

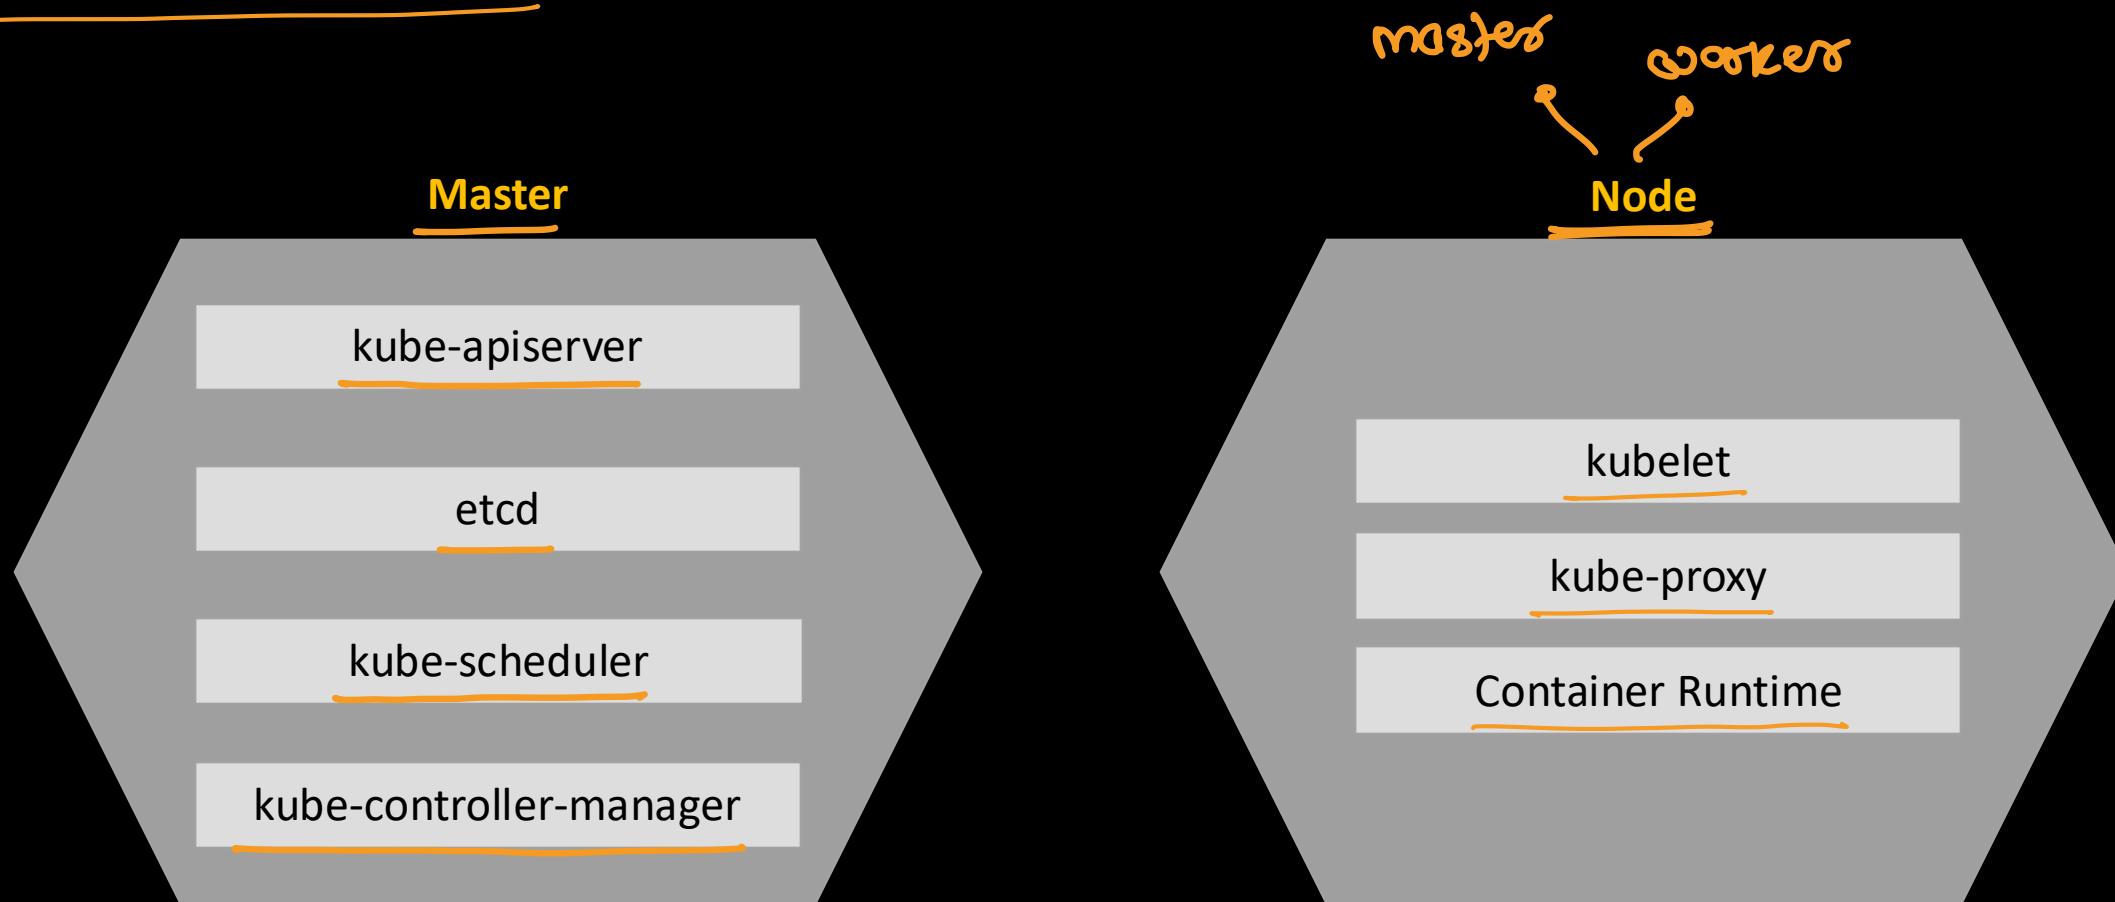
- one master & one or more workers
- used in dev or testing env

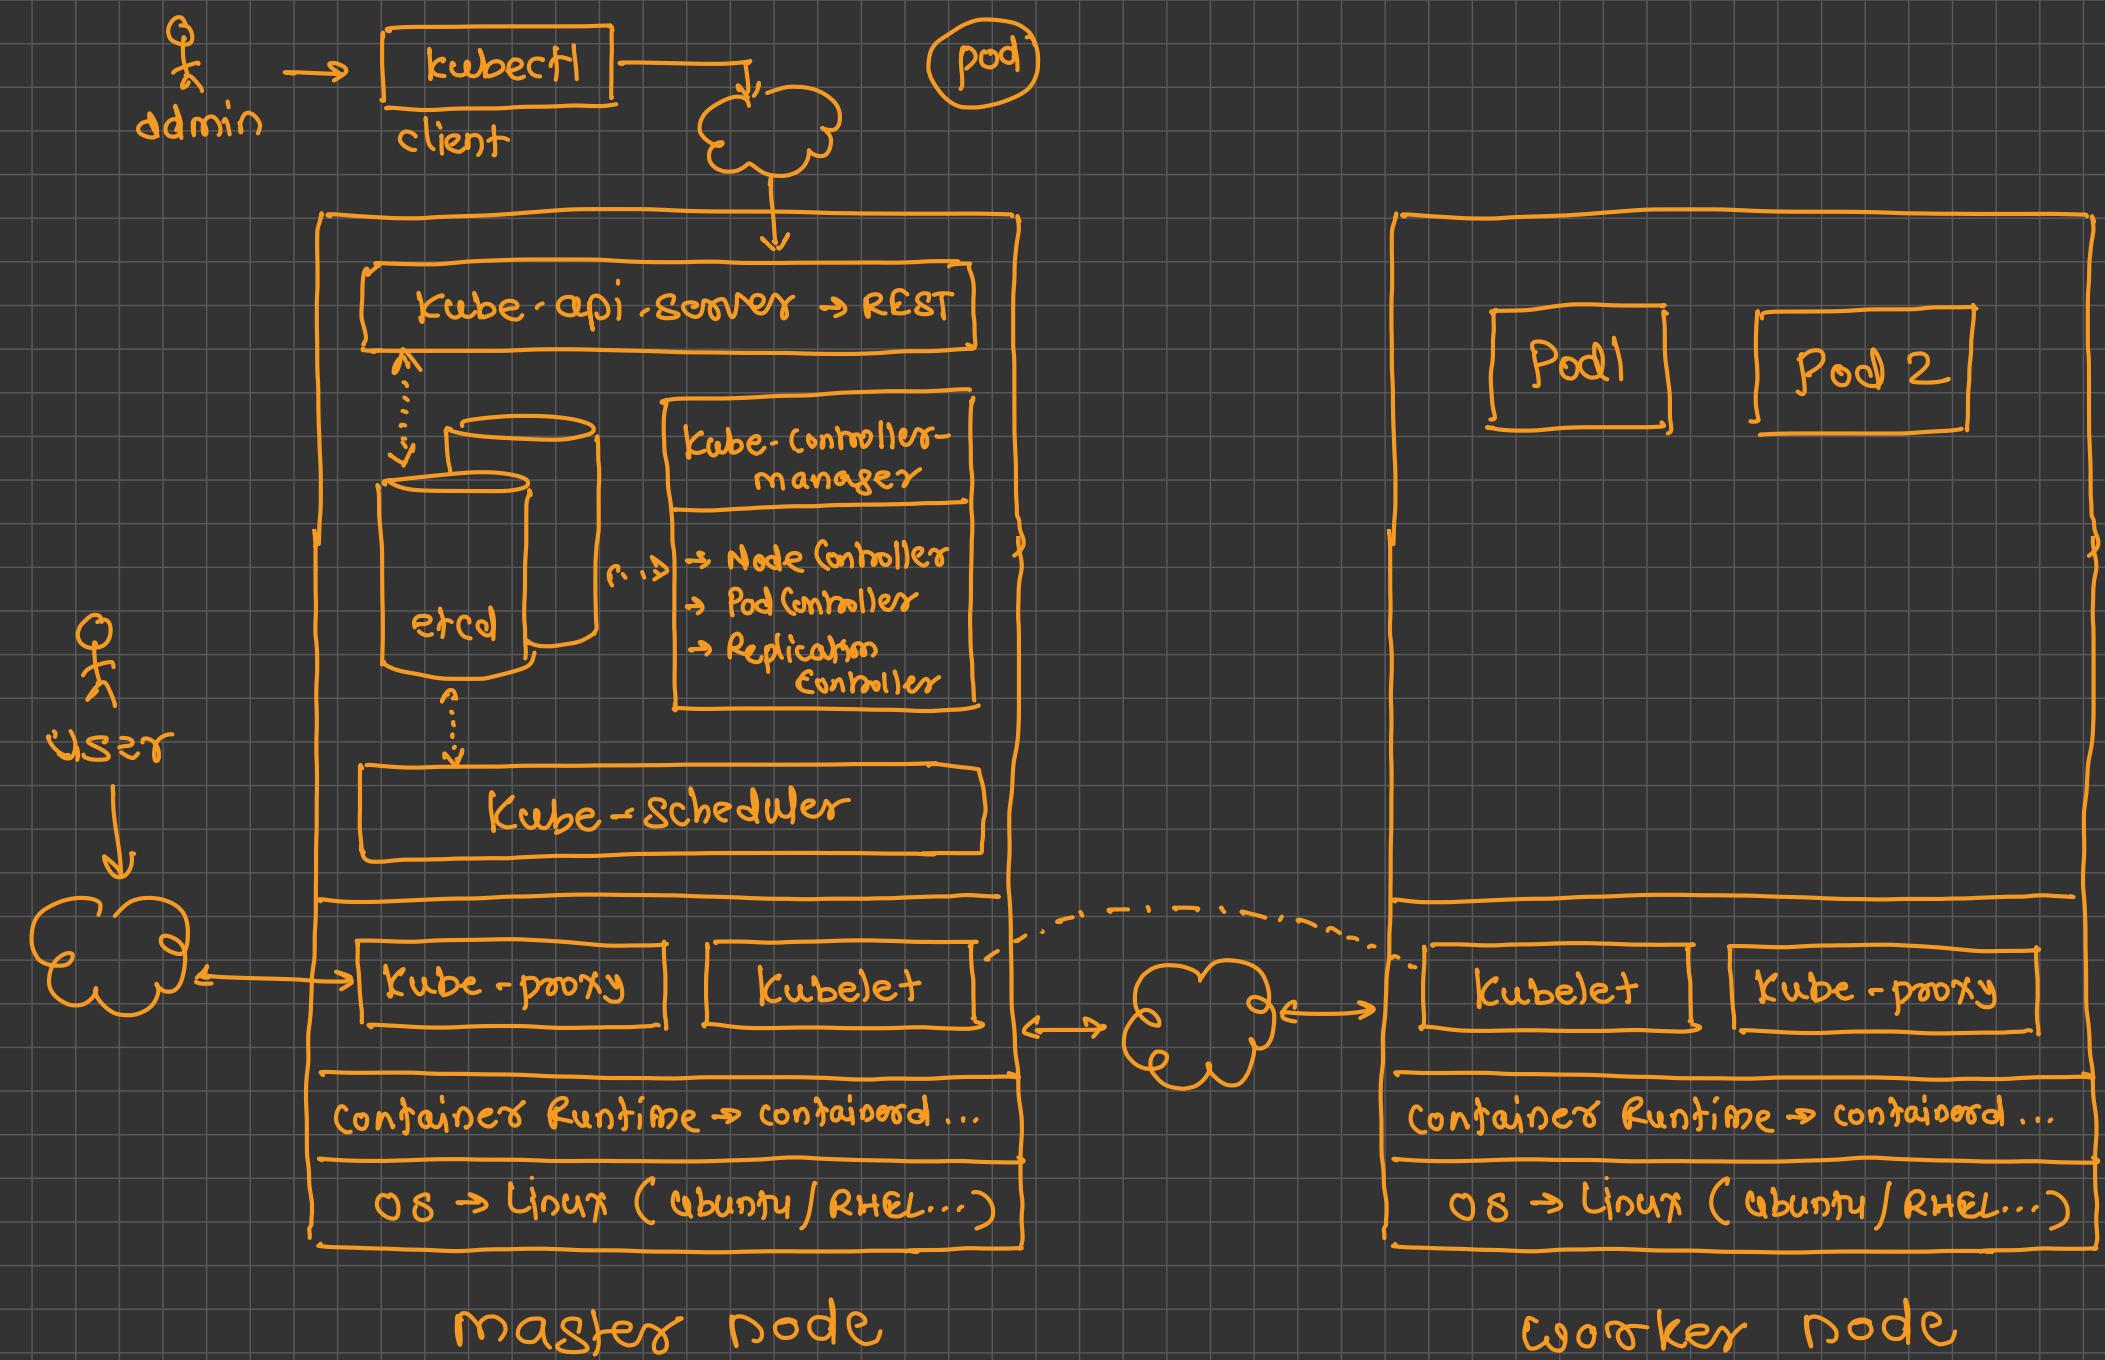
→ multi-master cluster

- multiple masters with multiple workers
- Highly Available cluster → HA
- preferred in production env.
- from multiple masters only one will be elected as leader
- K8S uses Raft consensus algo



Kubernetes Components







Master Components

- These are the **brain** of the Kubernetes cluster — they make all the decisions about what runs where, how scaling happens, and how the system stays healthy.
- Master components can be run on any machine in the cluster

Component	Function
<u>kube-apiserver</u>	The front door of the Kubernetes control plane. It exposes the Kubernetes API , which is used by all components, <code>kubectl</code> , and external clients.
<u>etcd</u>	The database of Kubernetes — stores all cluster data, configurations, and states in a key-value format.
<u>kube-scheduler</u>	Decides which node each Pod should run on based on resource requirements and constraints.
<u>kube-controller-manager</u>	Ensures that the actual state of the cluster matches the desired state defined in manifests.
<u>cloud-controller-manager</u> (optional)	Manages integration between Kubernetes and your cloud provider (e.g., AWS, GCP, Azure).



Master Components

■ kube-apiserver → brain → exposes REST APIs → consumed by kubectl command

- Acts as the communication hub between users, components, and the cluster.
- Every kubectl command goes through it.
- Validates requests and updates etcd accordingly.
- It's a stateless service – you can run multiple instances for high availability → multiple masters

■ etcd

- A distributed, consistent key-value store. → database
- Stores all cluster data, including:
 - Pod states
 - Configurations
 - Secrets
 - Node information
- It's the source of truth for your cluster.
- It's critical – if etcd is lost, your cluster loses its state.



Master Components

▪ kube-scheduler ↗ created by Pod Controller

- Watches for new Pods that don't have a Node assigned.
- Chooses the best Node to run the Pod based on:
 - Resource requests (CPU, memory)
 - Node affinity/anti-affinity
 - Taints and tolerations
 - Pod priorities
 - Custom policies

▪ kube-controller-manager

- Runs a set of controller loops, each responsible for maintaining part of the system's desired state.
- Examples of controllers:
 - Node Controller – manages node status.
 - Replication Controller – ensures the desired number of pod replicas are running.
 - Endpoint Controller – manages endpoint objects.
 - Service Account & Token Controllers – create default accounts and tokens.

↗ Replicaset
↗ DaemonSet
↗ StatefulSet

↳ Security



Node Components

→ master
→ worker

- Node components run on **every node**, maintaining running pods and providing the Kubernetes runtime environment

kubelet

- Primary agent that runs on every node.
- Communicates with the API Server.
- Ensures the containers defined in PodSpecs are running and healthy.
- Reports node and pod status back to the control plane. → Leaders
- Watches for Pod definitions assigned to the node and runs them using the container runtime.

kube-proxy

- Maintains network rules on nodes.
- Ensures that networking is consistent across the cluster. → overlay Network
- Implements Kubernetes Service abstraction – enabling communication between different Pods and Services.
- Uses iptables or IPVS to forward traffic to the correct Pod endpoints. → firewalls
- Supports load balancing between Pods behind a Service.

Container Runtime

- The actual software responsible for running containers.
- The kubelet interacts with the runtime through the Container Runtime Interface (CRI).
- Common runtimes:
 - containerd (default on most modern clusters)
 - CRI-O
 - Docker Engine (deprecated as of K8s v1.24+)
 - Mirantis Container Runtime



Create Cluster

- Use following commands on both master and worker nodes

```
> sudo apt-get update && sudo apt-get install -y apt-transport-https curl  
→ > curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -  
> cat <<EOF | sudo tee /etc/apt/sources.list.d/kubernetes.list deb https://apt.kubernetes.io/kubernetes-xenial main EOF  
> sudo apt-get update  
> sudo apt-get install -y kubelet kubeadm kubectl  
> sudo apt-mark hold kubelet kubeadm kubectl
```



Initialize Cluster Master Node

- Execute following commands on master node

```
> kubeadm init --apiserver-advertise-address=<ip-address> --pod-network-cidr=10.244.0.0/16  
> mkdir -p $HOME/.kube  
> sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
> sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

- Install pod network add-on

```
> kubectl apply -f  
https://raw.githubusercontent.com/coreos/flannel/2140ac876ef134e0ed5af15c65e414cf26827915/Documentation/kube-flannel.yml
```



Add worker nodes

- Execute following command on every worker node

```
> kubeadm join --token <token> <control-plane-host>:<control-plane-port> --discovery-token-ca-cert-hash sha256:<hash>
```



Steps to install Kubernetes

5

initialize

4

POD Network

3

kubeadm

2

docker

1

join

join

Master

Worker Node 1

Worker Node 2





Kubernetes Objects

- The basic Kubernetes objects include

- Pod
- Service
- Volume
- Namespace

- Kubernetes also contains higher-level abstractions build upon the basic objects

- Deployment
- DaemonSet
- StatefulSet
- ReplicaSet
- Job



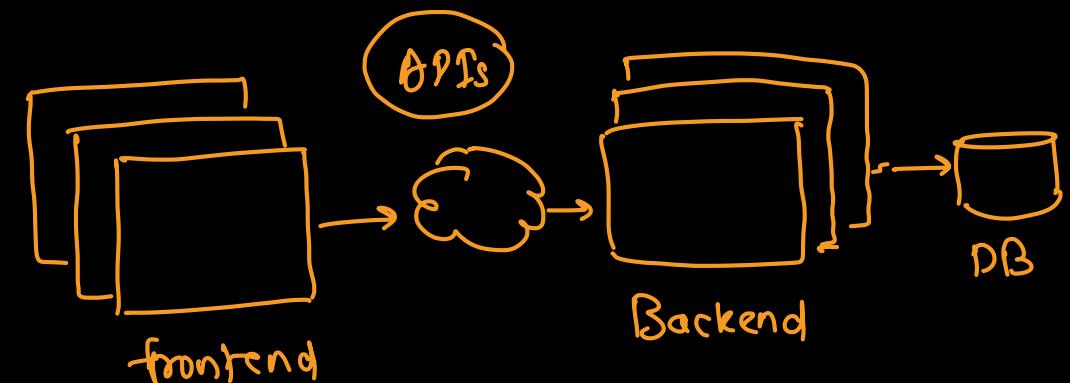
Namespace

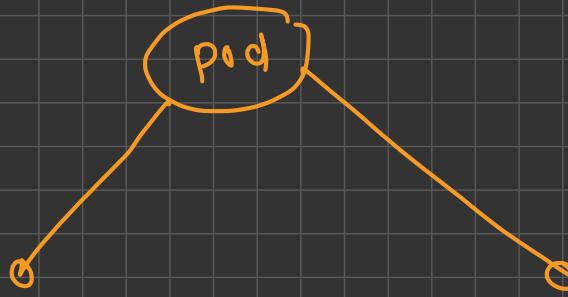
- Namespaces are intended for use in environments with many users spread across multiple teams, or projects
- Namespaces provide a scope for names
- Names of resources need to be unique within a namespace, but not across namespaces
- Namespaces can not be nested inside one another and each Kubernetes resource can only be in one namespace
- Namespaces are a way to divide cluster resources between multiple users



Pod

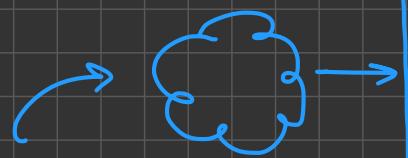
- A Pod is the basic execution unit of a Kubernetes application
 - The smallest and simplest unit in the Kubernetes object model that you create or deploy
 - A Pod represents processes running on your Cluster
 - Pod represents a unit of deployment
 - A Pod encapsulates
 - application's container (or, in some cases, multiple containers)
 - storage resources
 - a unique network IP
 - options that govern how the container(s) should run
- ↳ configuration





single container pod

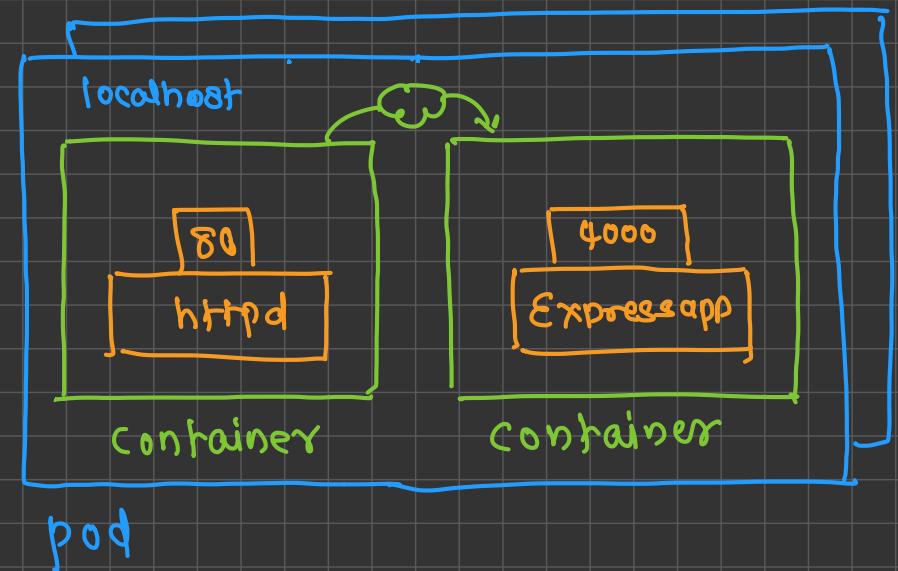
10-244.0.10



curl 10-244.0.10

multi-container pod

10.0.244.52



YAML to create Pod

apiVersion: v1

kind: Pod

metadata:

name: myapp-pod

labels:

app: myapp

spec:

containers:

- name: myapp-container

image: httpd

YAML for k8s objects

① apiVersion →

- version of k8s APIs

→ basic object → v1

higher level object = apps/v1

② kind →

- type of object to be created
- e.g. Pod, Service, Deployment

③ metadata →

- data/info about the object

- e.g. name, labels, namespace

④ spec →

- specification/definition of object

YAML → Syntax used to write config files

→ scalar → 10, person

→ map (key-value pair) → $\frac{\text{name}}{K} : \frac{\text{person}}{V}$

→ list (collection) →

- cat

- dog

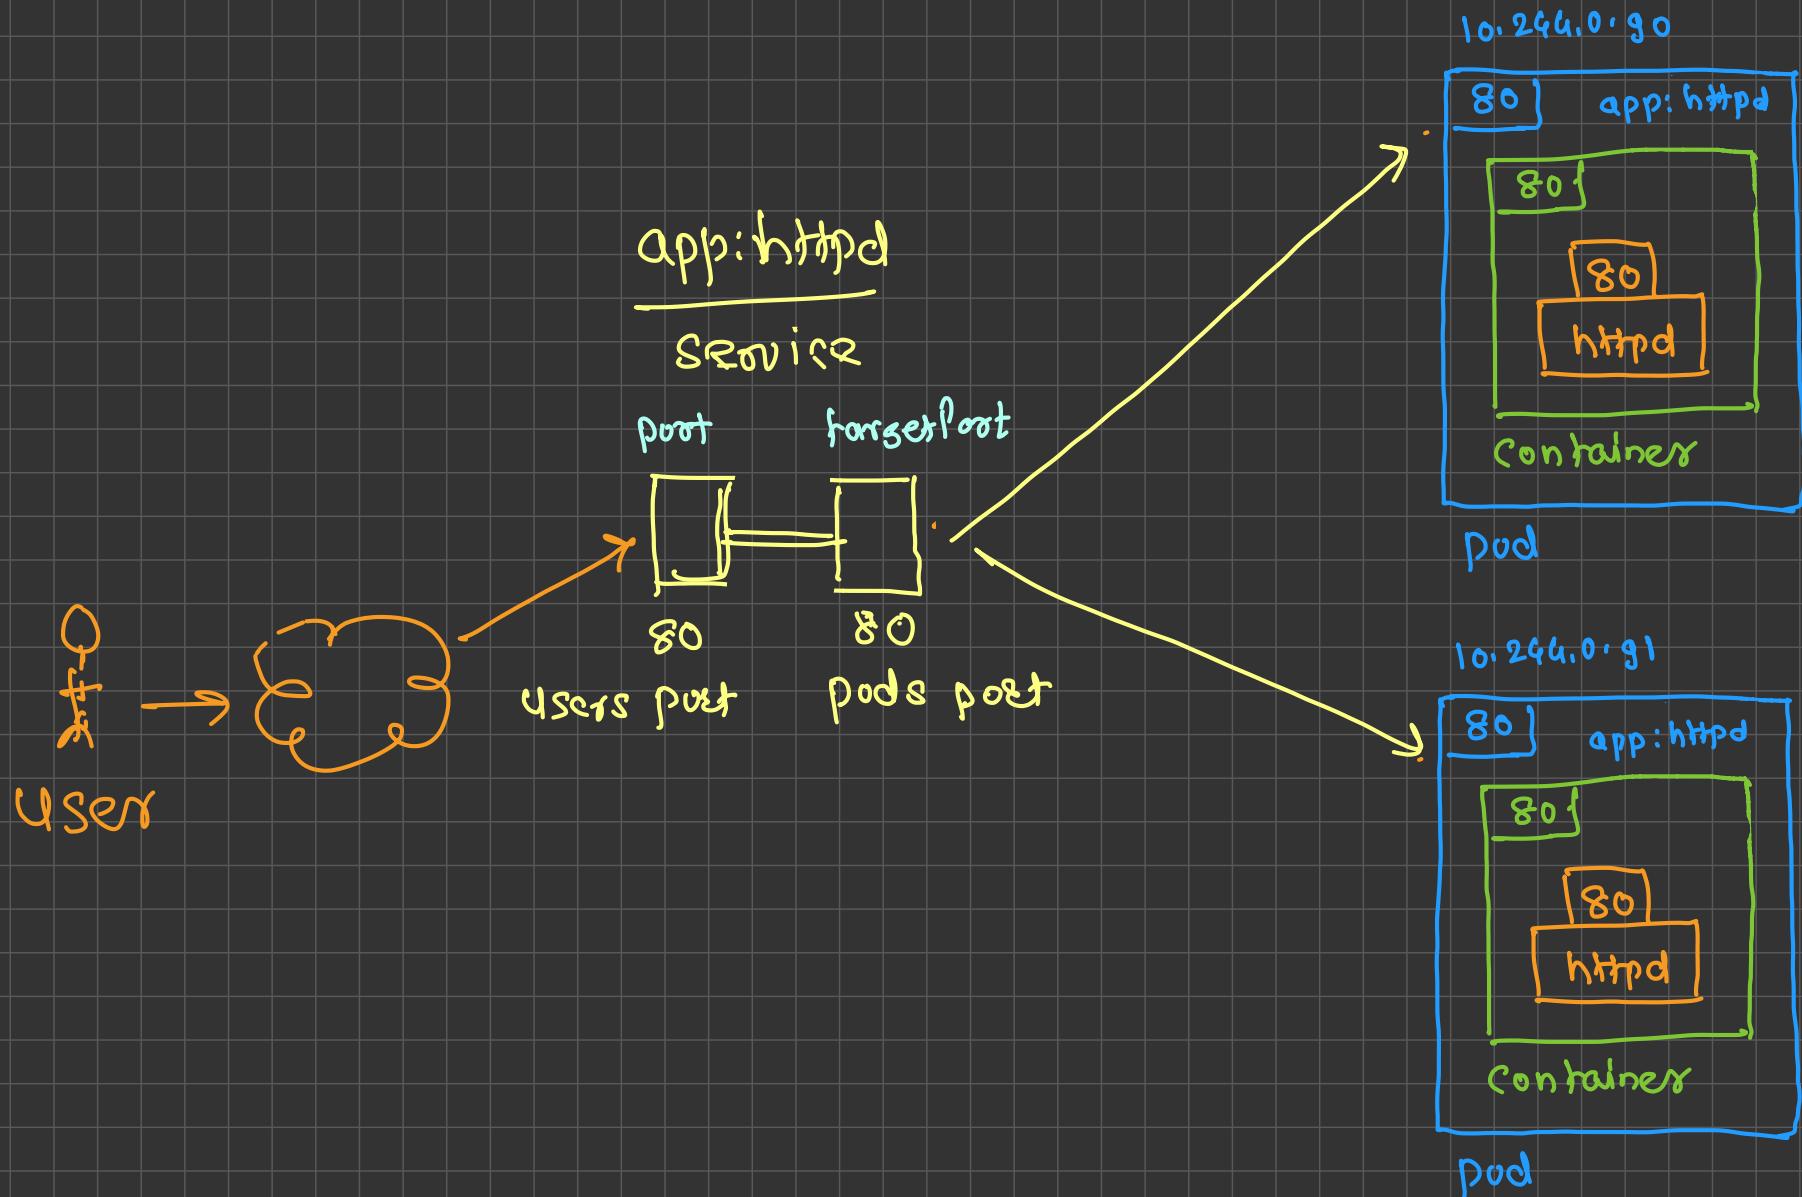
- horse



Service → load balancer

- An abstract way to expose an application running on a set of Pods as a network service
- Service is an abstraction which defines a logical set of Pods and a policy by which to access them (sometimes this pattern is called a micro-service)
- Service Types
 - ClusterIP
 - Exposes the Service on a cluster-internal IP
 - Choosing this value makes the Service only reachable from within the cluster
 - LoadBalancer
 - Used for load balancing the containers
 - NodePort

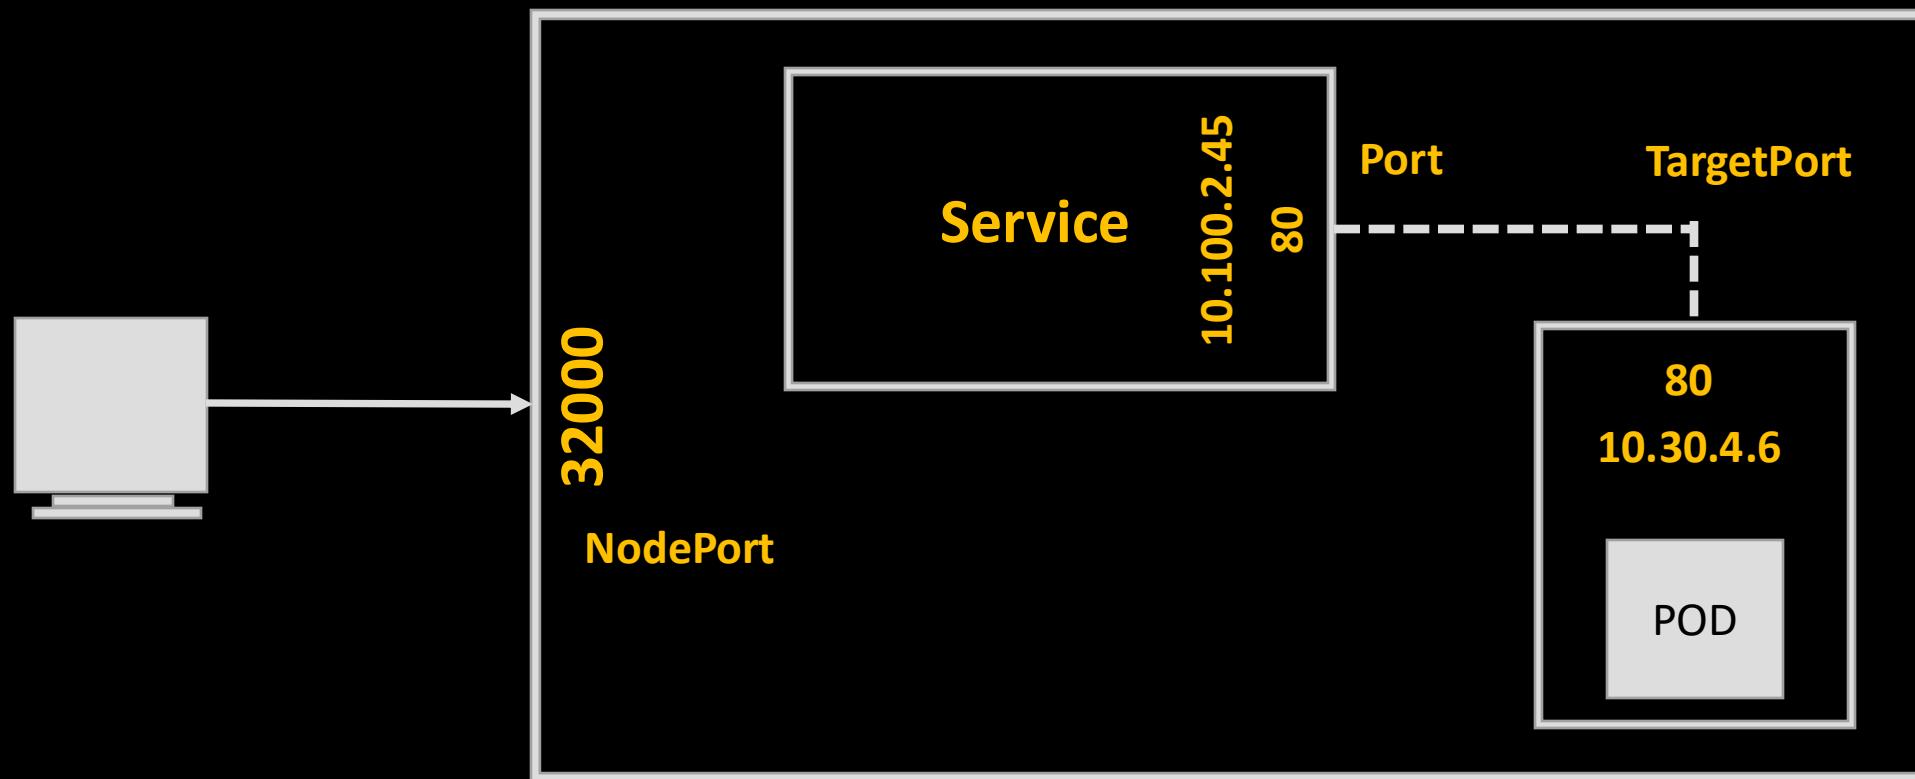
```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```





Service Type: NodePort

- Exposes the Service on each Node's IP at a static port (the NodePort)
- You'll be able to contact the NodePort Service, from outside the cluster, by requesting <NodeIP>:<NodePort>

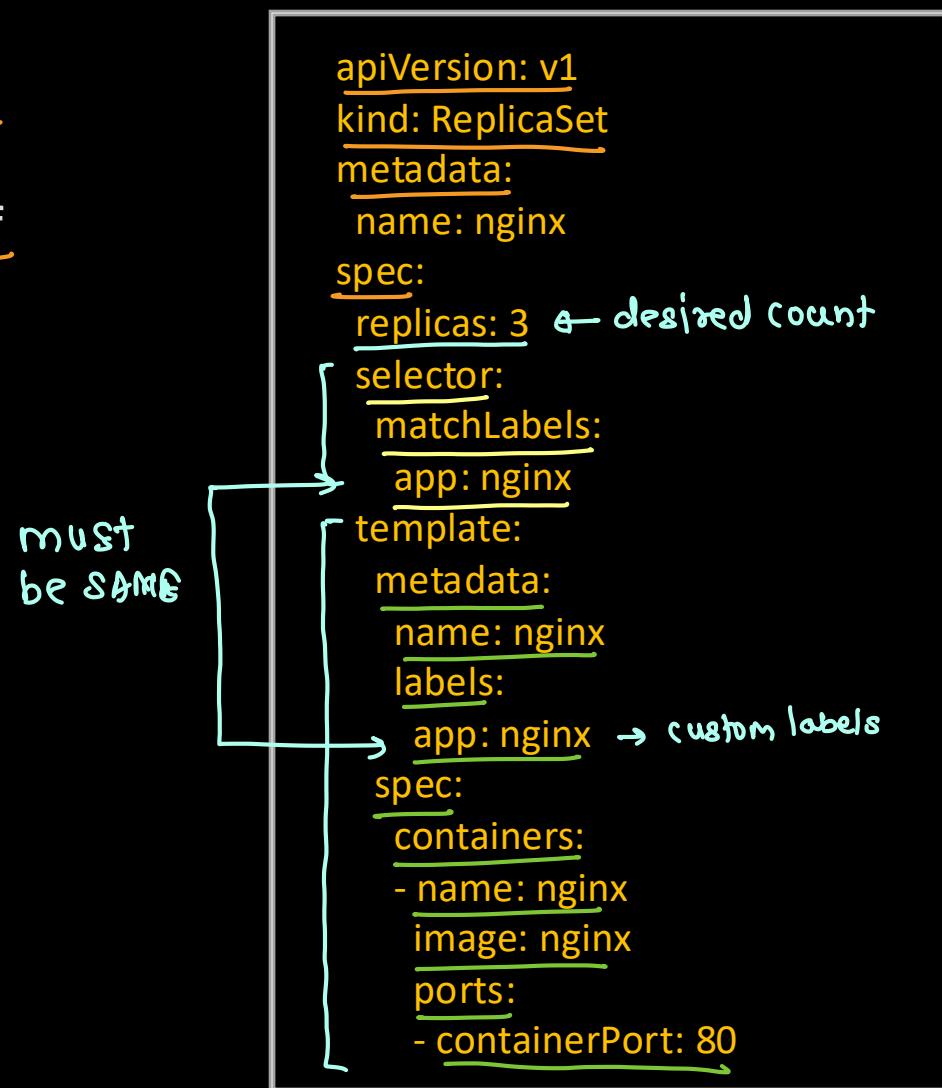
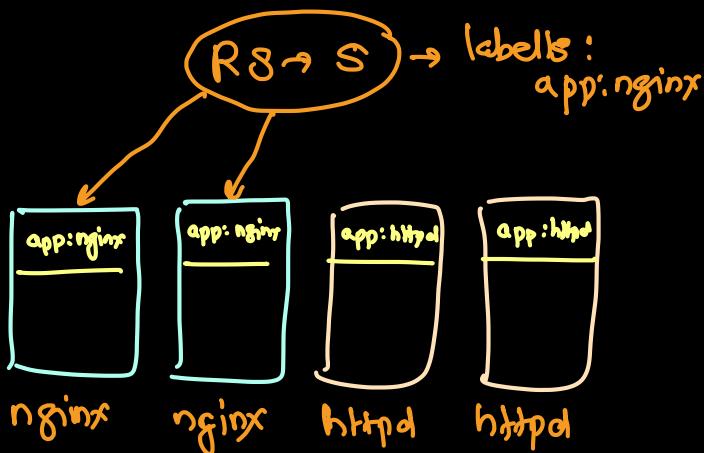




Replica Set

desired count

- A Replica Set ensures that a specified number of pod replicas are running at any one time
- In other words, a Replica Set makes sure that a pod or a homogeneous set of pods is always up and available
- If there are too many pods, the Replica Set terminates the extra pods
- If there are too few, the Replica Set starts more pods
- Unlike manually created pods, the pods maintained by a Replica Set are automatically replaced if they fail, are deleted, or are terminated





Deployment

- A Deployment provides declarative updates for Pods and ReplicaSets
- You describe a *desired state* in a Deployment, and the Deployment Controller changes the actual state to the desired state at a controlled rate
- You can use deployment for
 - Rolling out ReplicaSet
 - Declaring new state of Pods
 - Rolling back to earlier deployment version
 - Scaling up deployment policies
 - Cleaning up existing ReplicaSet

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: website-deployment
spec:
  selector:
    matchLabels:
      app: website
  replicas: 10
  template:
    metadata:
      name: website-pod
      labels:
        app: website
    spec:
      containers:
        - name: website-container
          image: pythoncpp/test_website
      ports:
        - containerPort: 80
```